

Kwame Osei

Data Engineer | Batch and streaming pipelines

Manchester · kwame.osei@correo-ficticio.uk · +44 7700 000006

PROFILE

Data engineer with five years of experience building reliable data pipelines. I focus on data modelling, orchestration and making analytics datasets trustworthy enough that downstream teams stop maintaining their own shadow copies. I care more about clear contracts between producers and consumers than about any particular tool.

EXPERIENCE

Data Engineer at Monzo (2022-2026). Built and maintained batch pipelines in Airflow feeding the analytics warehouse, covering roughly two hundred daily tasks. Designed dimensional models in dbt and introduced a layered structure separating raw ingestion from business logic. Introduced data quality tests and freshness checks that cut data incident reports by half in two quarters. Led the migration from a nightly full reload to incremental models, taking the warehouse refresh from four hours to twenty minutes. Junior Data Engineer at Deliveroo (2021-2022). Ingestion jobs in Python and Spark for order and rider event streams, plus backfills for historical data. Analytics intern at the Co-operative Group (2020-2021).

PROJECTS

Streaming pipeline consuming Kafka topics into near real time tables for the operations dashboard, replacing a fifteen minute polling job. Internal catalogue of dataset ownership and service level expectations, which made it possible to route data incidents to the right team automatically.

TECHNOLOGIES

Python, SQL, Apache Airflow, dbt, Apache Spark, Kafka, Snowflake, BigQuery, Postgres, Docker, Terraform, Git, Great Expectations, Looker.

EDUCATION

BSc Computer Science, University of Manchester (2017-2020).

ACHIEVEMENTS

Cut the monthly warehouse compute bill by 31% by rewriting the three most expensive models and partitioning the largest fact table. Wrote the team handbook on data modelling conventions, which is now the reference used in code review and onboarding. Reduced onboarding time for new analysts from three weeks to five days by documenting the core datasets and their known caveats.

ADDITIONAL TRAINING

Advanced dbt modelling workshop, 2024. Course on streaming systems and exactly once semantics, 2023. Internal training on data privacy and handling of personal data under GDPR, 2022.

OTHER

Available at one month notice. Open to hybrid work in Manchester or fully remote within the United Kingdom or Spain. Maintainer of a small open source package for Airflow sensors, with around four hundred monthly downloads. Occasional speaker at the Manchester data engineering meetup.

LANGUAGES

English native. Spanish B1. Twi native.